When Does A* Save Work? Controlled Experiments on Grids and Road Networks

Elnur Haqqverdiyev and Rajab Iskandarli
Department of Data Science and Artificial Intelligence
UFAZ, Advanced Data Structures and Algorithms, Fall 2026

***Abstract*—A* is often described as Dijkstra's algorithm with a heuristic, but the heuristic is not free: every generated vertex needs an estimate, and the heap entry becomes larger. We ask when the reduction in expanded vertices is large enough to pay for this extra work. We implement Dijkstra and A* in C++20 with the same CSR graph, binary heap, lazy deletion and search loop, and compare them on five unit-cost grid families (open, random obstacles, clustered obstacles, rooms and perfect maze) and on the SNAP California road network with an ALT landmark heuristic. The protocol is paired, pre-registered and frozen: 330 query instances, 3,600 timed runs, independent BFS distance oracles, and median/IQR plus bootstrap and Wilcoxon analysis. A* is dramatically better on open grids (105.7x median time, 222.4x fewer expansions), clearly better with random, clustered or room obstacles (12.8x, 12.8x and 8.7x time), but *slower* in mazes (0.81x, i.e. 24% slower) because Manhattan distance does not see the forced turns. On the road graph, ALT with 8 landmarks gives 7.62x speedup and 11.54x fewer expansions; 16 landmarks improve further to 13.52x and 22.41x. Time savings are systematically smaller than expansion savings because A* costs about 1.2–2.0x more per settled vertex. No distance mismatch or timeout occurred. The result is a practical rule rather than a universal claim: A* pays when the heuristic is informative about the actual remaining path.**

***Index Terms*—shortest paths, A*, Dijkstra, binary heap, admissible heuristic, ALT, road networks, experimental algorithmics, reproducibility.**

I. Introduction

Shortest-path queries are a basic primitive in routing, network analysis and AI search. Dijkstra's algorithm [1] is the standard exact method for non-negative edge weights; A* [2] uses an estimate of the remaining distance to focus the search. The textbook intuition is that A* expands fewer vertices and is therefore faster. In practice, the estimate is evaluated at every generated vertex, its values enter the heap key, and a weak heuristic can leave the explored region almost unchanged. The benefit is therefore conditional on both graph structure and heuristic quality.

This paper asks the T01 question: *how do graph structure and heuristic strength change the number of expanded vertices and the runtime of A* compared with Dijkstra, on unit-cost grids with obstacles and on a real-life graph?* We answer with a controlled experiment rather than a single benchmark. The main contributions are:

- a C++20 implementation with a documented CSR graph, binary heap, grid generators, Dijkstra, A*, Manhattan and ALT heuristics, and independent BFS/Floyd–Warshall oracles;
- a pre-registered protocol that separates expansion savings from time savings and pairs every A* run with the same Dijkstra query;
- a structure–heuristic map of when A* wins, including a clear negative result in mazes and a positive result on a real road network;
- an accounting of the per-expansion overhead that explains why time savings are smaller than expansion savings.

Section II reviews the algorithms. Section III describes the graphs and Section IV the implementation and protocol. Section V reports the measurements; Sections VI and VII discuss the meaning and limits.

II. Background

A. Dijkstra and A*

Let $G = (V, E, w)$ be a graph with $w(u, v) > 0$. Dijkstra maintains a tentative distance $g(v)$ and repeatedly settles the unsettled vertex with minimum g. With a binary heap and lazy deletion the complexity is $O((n + m) \log n)$ in the worst case, where $n = |V|$ and $m = |E|$. A* additionally uses a heuristic $h(v)$ estimating the distance from v to the target t, and orders the open list by

$$f(v) = g(v) + h(v). \tag{1}$$

If h is admissible, $h(v) \le d(v, t)$, A* returns a shortest path; if it is consistent, $h(u) \le w(u, v) + h(v)$, settled vertices never need reopening [2], [4]. Both heuristics used here are consistent.

B. ALT landmarks

For an undirected graph and exact distances $d(L, v)$ from a landmark L, the triangle inequality gives

$$|d(L, v) - d(L, t)| \le d(v, t). \tag{2}$$

The ALT heuristic takes the maximum over a set $\mathcal{L}$ of landmarks [3]:

$$h(v) = \max_{L \in \mathcal{L}} |d(L, v) - d(L, t)|. \tag{3}$$

It is admissible because every term is a lower bound, and consistent because each term changes by at most the edge weight across an edge. The landmark distance tables are computed by BFS before any timed search.

TABLE I
EXPERIMENTAL GRAPHS AND QUERY BUDGET.

Data	graphs	queries	repeats	timed runs
Grid controls	20	300	3	2,700
Road network	1	30	5	900
Total	21	330	–	3,600

C. Expansions versus time

Let X be the number of settled vertices. If the average cost of an expansion is c_D for Dijkstra and c_A for A*, then

$$\frac{T_D}{T_A} = \frac{X_D\, c_D}{X_A\, c_A}. \tag{4}$$

The expansion ratio X_D/X_A overstates the time ratio whenever $c_A > c_D$. This decomposition is the main explanatory tool of the paper.

III. DATA

A. Grid controls

Grids are square, row-major and 4-neighbour with unit edge cost. A blocked cell is an isolated vertex; only the largest 4-connected component is used for query selection. We use five families:

- **open**: no blocked cells;
- **random**: independent blocking with probability 0.20;
- **clustered**: 30% blocked in random circular blobs;
- **rooms**: walls every 16 cells with one door per room-boundary segment;
- **maze**: a perfect maze generated by iterative randomised DFS.

Sizes are 128^2, 256^2, 512^2 and 1024^2 cells. These families isolate different structural effects: exact geometric guidance (open), local detours (random), long detours (clustered/rooms) and forced corridors (maze).

B. Real road network

We use SNAP `roadNet-CA`, the California road network [5]. The SNAP page reports 1,965,206 nodes and 2,766,607 edges. We read the 5,533,214 directed arcs, symmetrise each road segment, remove self-loops and duplicates, take the largest connected component and compact the vertex ids. The resulting graph has **1,957,027 vertices** and **2,760,388 undirected edges** (5,520,776 CSR arcs), average degree ≈ 2.82. SNAP does not provide road lengths or travel times, so we use **unit hop cost** rather than inventing a weighted interpretation. The result is therefore a hop-count result; it is not a claim about travel-time routing.

IV. IMPLEMENTATION AND PROTOCOL

A. Graph, heap and search

The graph is stored in compressed sparse row (CSR) form: `offsets[n+1]`, `targets[m]` and `weights[m]` use $O(n+m)$ memory, degree is $O(1)$ and neighbour iteration is $O(\deg u)$.

The binary heap is an array-based min-heap with duplicate entries and lazy deletion. `push` and `pop` are $O(\log n)$ worst case; `top` is $O(1)$. There is no decrease-key: a popped entry is stale if its g differs from the current tentative distance or if the vertex is already closed. Two deterministic tie-breaks are compared: `node_id` uses (f, id); `larger_g` uses $(f, -g, id)$.

Both algorithms call the *same* `shortest_path` routine. A null heuristic is Dijkstra; a non-null heuristic is A*. This guarantees the same CSR traversal, heap, relaxation and stopping rule. The search stops when the target is popped and settled. Timed runs do not reconstruct paths; correctness runs do. Scratch distance/closed/parent arrays are reused and reset in $O(1)$ with a generation stamp, so allocation is outside the timed region.

B. Heuristics

On grids we use Manhattan distance

$$h(v) = |row(v) - row(t)| + |col(v) - col(t)|. \tag{5}$$

It is admissible and consistent for 4-neighbour unit grids, with or without obstacles.

On the road graph we use ALT (3) with 16 landmarks selected by deterministic farthest-point sampling. The first landmark is a fixed seed; each next landmark maximises the minimum BFS distance to the already selected landmarks. The main treatment uses $K = 8$; $K = 1, 4, 16$ are secondary treatments.

C. Correctness first

Before timing, the test suite checks: heap order against a reference priority queue; a hand-computed weighted graph against Floyd–Warshall; a hand-checked 5×5 wall grid with distance 8; a no-path grid; all five grid families against independent BFS; ALT admissibility and consistency for every free vertex; tie-break distance equivalence; and a sample of real-graph queries with edge-by-edge path validation.

Every benchmark query is also checked against BFS before it is recorded. A distance mismatch aborts the run. Timeouts are recorded as `timeout`, never silently dropped. All 3,600 main runs returned `ok` and matched the oracle; no timeout occurred.

D. Queries and timing

Queries are selected by BFS only, independently of Dijkstra/A*. For a random source, BFS gives eccentricity $d_{\max}$; targets are sampled in three bands: short $[0.08, 0.25]d_{\max}$, medium $[0.40, 0.60]d_{\max}$ and long $[0.80, 0.97]d_{\max}$. Each grid has 5 sources (15 queries); the road graph has 10 sources (30 queries).

Each treatment gets one untimed warm-up per graph. Grid queries are repeated three times and road queries five times; the median is the query-level summary and the raw repeats are kept. Dijkstra and A* run on the same query and their order alternates. The timed region contains only the search call; graph construction, BFS query selection and landmark

TABLE II
ADT INTERFACES AND WORST-CASE COSTS (n VERTICES, m ARCS).

Component	Interface and cost
CSR graph	neighbour loop $O(\deg u)$; degree $O(1)$; memory $O(n+m)$
Binary heap	push/pop $O(\log n)$; top $O(1)$; lazy deletion
Search	Dijkstra $h \equiv 0$; A* $f = g + h$; same loop
Manhattan	$O(1)$ per estimate; admissible/consistent
ALT	$O(K)$ per estimate; K landmark tables
Oracle	BFS $O(n+m)$; Floyd–Warshall $O(n^3)$ for tiny tests

tables are outside it. A 30 s per-search cap is checked every 4096 pops and a 4 GiB process cap is applied externally. Runs are pinned to one CPU with `taskset`.

E. Counters and statistics

Expanded is a popped entry that is not stale and not already closed (the target counts when popped). *Pushed* counts every heap insertion, including the source and every successful decrease. *Stale_popped* counts popped entries whose g is outdated or whose vertex is closed. *Relaxed_edges* counts outgoing arcs inspected from expanded vertices. *Heuristic_calls* counts heuristic estimates (zero for Dijkstra).

We report medians and interquartile ranges over query-level medians. Confidence intervals for medians use 10,000 paired bootstrap resamples. Paired Wilcoxon signed-rank tests are reported for time and expansions. Within-query repeat variation is reported as a coefficient of variation (CV); the median CV was 7.2% for grid Dijkstra, 10.1% for grid A*, 7.4% for road Dijkstra and 9.6% for road A* ($K = 8$). Expansions were identical across repeats.

V. RESULTS

A. Grid families

Table III gives the main treatment (Manhattan, `larger_g`). Ratios are Dijkstra/A*; values above 1 mean A* wins. The open grid is the best case: A* expands 222.4x fewer vertices and is 105.7x faster in the median, and it is faster on 100% of queries. Random and clustered obstacles still give a clear win (21.4x and 22.9x fewer expansions; 12.8x time), but the advantage is an order of magnitude smaller than on the open grid. Rooms are weaker still: 14.3x expansions and 8.7x time.

The maze is the counterexample. The expansion ratio is 1.02 (about 2% fewer expansions for A*) and the time ratio is 0.81, so A* is about 24% *slower*. Manhattan distance sees the straight line to the target but not the forced turns of the maze. A* pays the heuristic cost without changing the explored region.

The open-grid advantage grows with size: at 128^2 the median is 45.5x time and 94.4x expansions; at 1024^2 it is 389.9x and 696.3x. Random grids are roughly flat in ratio (12.0x at 128^2, 21.0x at 1024^2), while rooms do not improve monotonically (7.6x, 8.7x, 11.4x, 8.7x). The maze ratio rises from 0.75x to 0.97x as the fixed overhead becomes small relative to the search, but A* never becomes faster in the median.

TABLE III
GRID RESULTS, MANHATTAN HEURISTIC, `larger_g` TIE-BREAK. RATIO IS DIJKSTRA/A*; IQR IN PARENTHESES.

Family	Time ratio	Expansion ratio	A* faster
Open	105.7 (54.9–249.3)	222.4 (105.1–490.3)	100%
Random 20%	12.8 (7.2–21.9)	21.4 (11.3–31.7)	98%
Clustered 30%	12.8 (7.0–34.6)	22.9 (11.7–67.8)	98%
Rooms	8.7 (5.6–17.0)	14.3 (8.4–25.6)	100%
Maze	**0.81** (0.75–0.94)	**1.02** (1.00–1.05)	17%

TABLE IV
ROAD NETWORK RESULTS WITH ALT; RATIO IS DIJKSTRA/A*.

Treatment	Time ratio	Expansion ratio	Time saving	Faster
K=1	1.29 (0.93–1.75)	1.96 (1.44–2.44)	22.5%	70%
K=4	4.75 (2.38–7.71)	7.26 (3.76–10.86)	78.9%	100%
K=8	**7.62 (4.38–16.65)**	**11.54 (7.06–25.66)**	86.8%	100%
K=16	**13.52 (7.64–25.25)**	**22.41 (12.94–38.70)**	92.6%	100%

Distance bands are not a universal explanation. On the open grid, long queries benefit more than short ones (156.0x vs 83.5x time). On random grids the opposite ordering appears (long 11.5x vs short 14.4x): a long random query has more opportunity to accumulate detours that Manhattan cannot predict. Structure, not distance alone, controls the benefit.

B. Real road network

Table IV reports ALT. With one landmark the expansion saving is real (1.96x) but the time saving is not statistically significant (1.29x, $p = 0.13$): the heuristic is too weak to pay for itself. With 4, 8 and 16 landmarks both differences are significant ($p < 10^{-8}$). The main treatment, $K = 8$, is 7.62x faster and expands 11.54x fewer vertices; $K = 16$ reaches 13.52x and 22.41x.

The K=16 result rejects the expectation of a plateau at K=8: paired on the same 30 queries, K=16 uses about 21% less time and 25% fewer expansions than K=8 ($p < 10^{-4}$). More landmarks keep tightening the bound faster than their per-expansion cost grows on this road graph.

By band (K=8), the time ratios are 14.9x (short), 9.4x (medium) and 4.3x (long); expansion ratios are 19.1x, 13.5x and 7.5x. Absolute work still grows with distance: median long-query A* expands 247,065 vertices and takes 99 ms, versus 448 ms for Dijkstra. A long road route gives the landmark bound more chances to be loose, so the ratio is smaller even though the absolute saving is large.

C. Why time differs from expansions

Table V shows the measured cost per expanded vertex. A* costs 1.2–2.0x more per expansion. On the open grid the expansion ratio is 222.4x but the time ratio is only 105.7x; on random grids 21.4x becomes 12.8x. In the maze the cost ratio is 1.21x and the expansion ratio is 1.02x, so the time

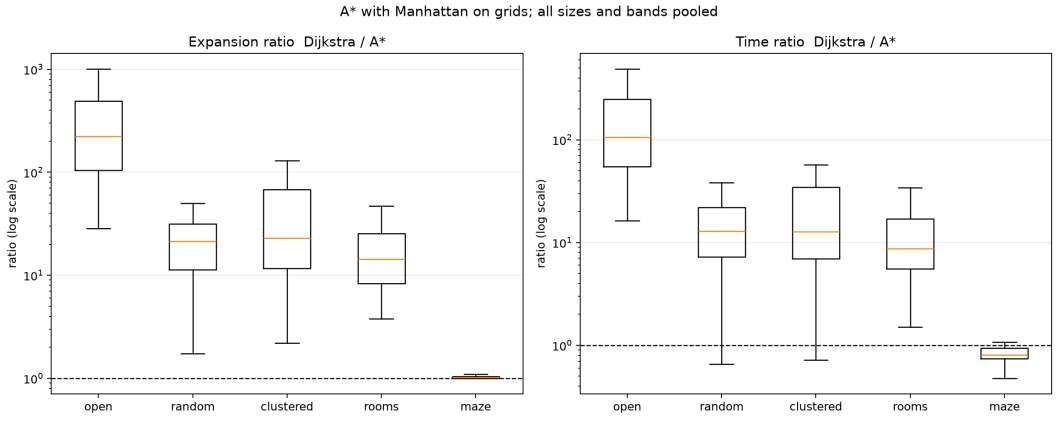


Fig. 1. Expansion and time ratios for the five grid families, all sizes and distance bands pooled. The horizontal line is parity. The maze is below parity in time even though its expansion ratio is near one.

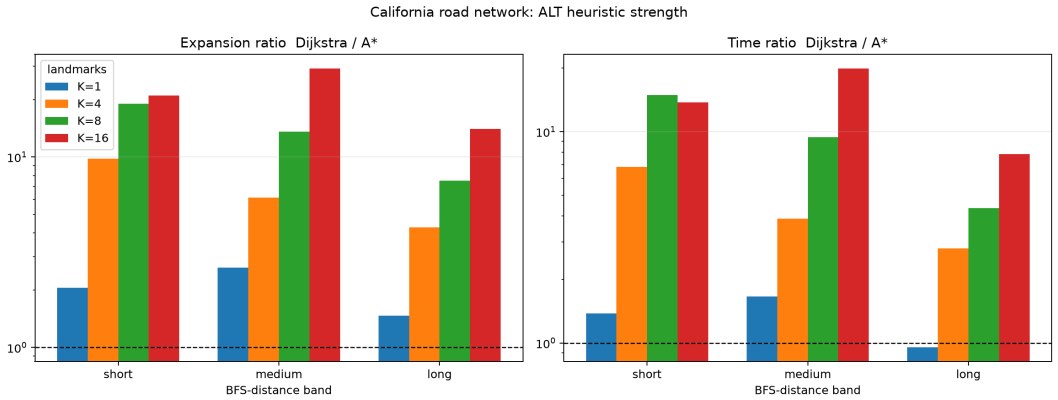


Fig. 2. Road network: expansion and time ratios by BFS-distance band and number of ALT landmarks. The y-axis is logarithmic; the dashed line is parity. The K=1 time ratio on long queries is below one.

TABLE V
COST PER EXPANDED VERTEX (NS), MEDIANS.

Data	Dijkstra	A*	Ratio
Open grid	142.2	288.3	2.03x
Random grid	156.5	254.1	1.62x
Clustered grid	126.7	245.0	1.93x
Rooms grid	146.0	270.4	1.85x
Maze	64.9	78.8	1.21x
Road, ALT K=8	234.5	372.9	1.59x

ratio is 0.81x: the extra work has no compensating reduction in expansions. On the road graph, Dijkstra costs 234.5 ns per expansion, ALT K=8 372.9 ns (1.59x) and K=16 393.7 ns (1.68x).

D. Tie-breaking

The `larger_g` rule has a median expansion ratio of 1.00 on the open grid relative to `node_id`, but `node_id` is worse on 45% of open-grid queries, up to 285.7x. The corresponding fractions are 63% (random), 68% (clustered), 58% (rooms) and 2% (maze). On open and clustered grids `larger_g` never

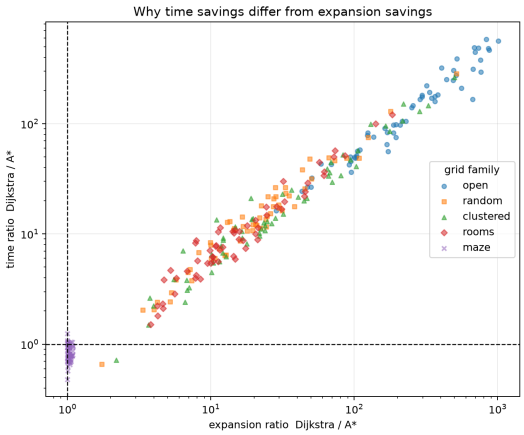


Fig. 3. Per-query expansion ratio versus time ratio. Points below the diagonal have time savings smaller than expansion savings; this is the expected effect of the per-expansion overhead.

expanded more vertices; on random grids the worst regression was 0.8%. Tie-breaking is therefore a tail-risk decision rather

TABLE VI
HEURISTIC DIAGNOSTICS AND OUTCOMES (GRID MEDIANS; EXPLORATORY).

Family	h corr.	norm. MAE	detour	time ratio	EBF A*
Open	1.00	0.00	1.00	105.7x	1.00
Random	1.00	0.02	1.00	12.8x	1.02
Clustered	0.99	0.04	1.00	12.8x	1.01
Rooms	0.99	0.06	1.02	8.7x	1.02
Maze	0.42	0.97	35.06	0.81x	1.00

than a universal median effect. It never changes the returned distance.

The frozen run also contains both tie-break rules on the real road graph with ALT K=8. There, `node_id` expands a median 1.14x as many vertices (+13.8%), is worse on 27 of 30 queries, and reaches 2.39x on the worst query (paired Wilcoxon $p = 4.4 \times 10^{-6}$). Median wall-clock time is 1.09x, but that difference is noisy ($p = 0.124$). The expansion count is the reliable signal, and it confirms the grid conclusion: `larger_g` is a justified default rather than a cosmetic choice.

E. Heuristic informativeness

The expansion result is descriptive; we now ask what makes a heuristic useful. For each grid query we run BFS from the target, sample 5,000 vertices, and compare $h(v)$ with the true distance $d(v, t)$. We report Spearman correlation, normalised mean absolute error $\mathrm{MAE}/\bar{d}$, detour ratio $d(s,t)/\mathrm{Manhattan}(s,t)$, and effective branching factor.

Table VI shows the expected ordering. The open grid has an exact heuristic: correlation 1.00, normalised MAE 0.00. Random, clustered and room grids retain a high rank correlation (0.99) but have growing error (0.02–0.06); their speedups fall from 105.7x to 8.7–12.8x. The maze has correlation 0.42 and normalised MAE 0.97; its speedup is 0.81x.

Across all 300 grid queries, the pooled Spearman correlation is $+0.78$ for h–distance correlation and -0.81 for normalised MAE. A single-feature least-squares model $\log(T_D/T_A) \sim \mathrm{MAE}/\bar{d}$ has $R^2 = 0.55$. This is an explanatory model, not a universal predictor: within one family the error range is narrow, and leave-one-family-out extrapolation to the maze is poor. The mechanism is nevertheless clear: admissibility is not enough; the heuristic must be *informative* about the actual remaining distance.

On the road graph (ALT K=8), the heuristic is also informative: h–distance correlation 0.98, normalised MAE 0.155. The query-level Spearman correlation between absolute heuristic error and time ratio is -0.77.

F. Price of optimality: weighted A*

The maze is the negative case for optimal A*. We therefore test the weighted objective $f = g + \epsilon h$ on all five grid families and the road graph, with $\epsilon \in \{1, 1.1, 1.25, 1.5, 2, 3, 5, 10, 20, 50\}$. The primary results use a closed list, where a settled vertex is not reopened; reopening is retained as a sensitivity check because it is the aggressive choice. All weighted runs returned at least the optimal distance; $\epsilon = 1$ was exact under both policies. Table VII shows selected rows.

A perfect maze has a unique simple path between two cells, so every returned path has the same cost. The measured path-cost ratio is exactly 1.00 for every ϵ, and at $\epsilon = 50$ weighted A* is 1.10x faster than Dijkstra. The optimality price is therefore zero on this control, but only because there is nothing to trade.

On graphs with alternative paths the trade is real and depends on the base heuristic. With the closed list, random obstacles at $\epsilon = 1.1$ improve the median speedup from 12.8x to 20.1x for a 1.8% median cost increase; $\epsilon = 1.25$ reaches 38.2x for 5.5%. Rooms at $\epsilon = 1.25$ reach 24.8x for 3.1%, and clustered obstacles at $\epsilon = 1.5$ reach 26.5x for 5.0%. The open grid is the control where Manhattan is exact: cost stays 1.00 and weighting cannot reduce expansions further.

The road graph exposes the policy choice. Optimal ALT K=8 is 7.4x faster than Dijkstra in this sweep; the closed-list $\epsilon = 1.1$ variant reaches 8.8x at 1.4% median cost, while reopening the same ϵ collapses to 5.4x. Where the heuristic is already tight, reheapifying settled vertices costs more than the shorter search saves. Weighting is thus useful when the heuristic is weak, not as a universal speed switch, and the closed list is the better default for this implementation.

G. Batch throughput with OpenMP

The problem statement asks about one query at a time, and a single A*/Dijkstra search is inherently sequential. Independent queries, however, are embarrassingly parallel. We run the five 1024^2 grid families and the 30 road queries (three repeats each) with 1, 2, 4, 8 and 16 OpenMP threads. Each task still runs Dijkstra and A* sequentially on the same query; each worker gets its own search scratch and heuristic object. Input construction, query selection and thread-local scratch allocation are outside the timed parallel region. The frozen query-level results remain serial.

Table VIII shows 9.73x batch throughput at 16 threads (61% parallel efficiency): 9.93x for the grid workload and 9.36x for the road workload. The sublinear scaling is expected from the mix of performance and efficiency cores and from memory bandwidth. The efficiencies slightly above one at two and four threads are within frequency/run-to-run variation, not super-linear scaling. Per-query latency inflates by 21% for Dijkstra and 18% for A* at 16 threads; the median Dijkstra/A* ratio rises from 14.8x to 17.9x, and the maze remains below parity (0.87x). The structural conclusion is therefore unchanged under parallel load. This is a throughput result, not a claim that a single query becomes faster.

VI. DISCUSSION

Table IX summarizes the pre-registered predictions. H1–H4 were confirmed. H5 was only partly confirmed: open-grid long queries benefit more than short ones, but random-grid long queries do not. H6 was partly rejected: K=8 is strong, K=1 does not produce a significant time saving, but K=16 is still better than K=8 rather than a plateau. H7 was rejected as a median claim; the tie-break matters in a heavy tail.

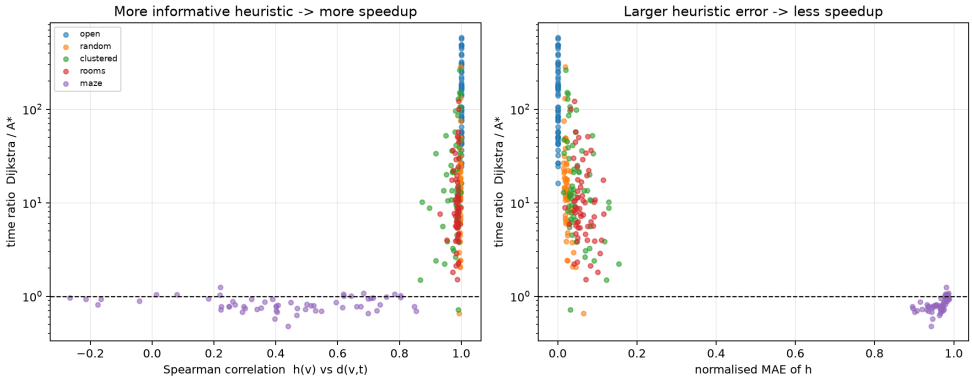


Fig. 4. Heuristic informativeness versus speedup. Left: rank correlation between $h(v)$ and the true remaining distance. Right: normalised MAE. The maze is the low-information, low-speedup outlier.

TABLE VII
WEIGHTED A* WITH A CLOSED LIST (EXCEPT THE LAST ROW). RATIO IS DIJKSTRA/WEIGHTED A*; COST IS RETURNED PATH COST DIVIDED BY THE OPTIMAL COST. SELECTED ϵ VALUES; THE FULL SWEEP HAS TEN VALUES.

workload	policy	ϵ	time ratio	cost median	cost p95	cost max
random	closed	1.10	20.1x	1.018	1.06	1.07
random	closed	1.25	38.2x	1.055	1.13	1.15
rooms	closed	1.25	24.8x	1.031	1.09	1.11
clustered	closed	1.50	26.5x	1.050	1.15	1.21
maze	closed	50.0	1.10x	1.000	1.00	1.00
road	closed	1.00	7.4x	1.000	1.00	1.00
road	closed	1.10	8.8x	1.014	1.04	1.04
road	reopen	1.10	5.4x	1.011	1.04	1.04

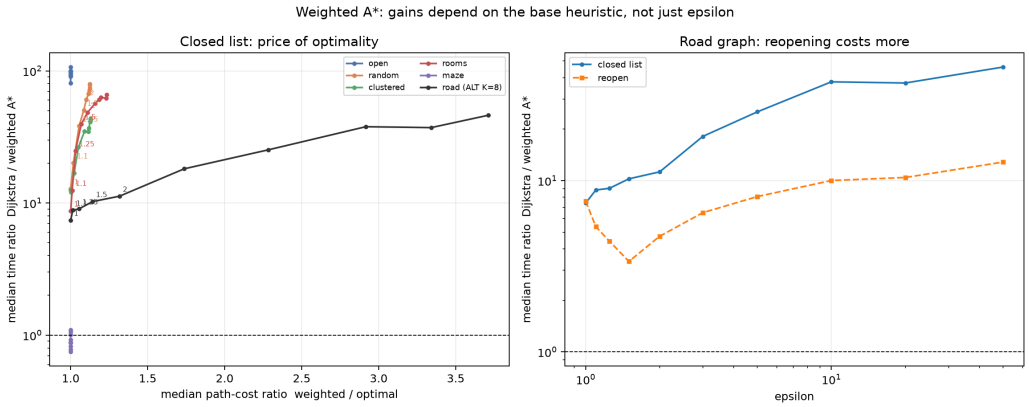


Fig. 5. Weighted A* on graphs with alternative paths. Left: closed-list cost versus speed, with selected ϵ labels. Right: on the road graph the closed list is faster at every $\epsilon > 1$, while reopening churns the heap for no useful gain. The open grid and perfect maze sit at cost ratio 1.00 because they leave no alternative path to trade.

TABLE VIII
OPENMP BATCH THROUGHPUT (FIVE 1024^2 GRIDS PLUS THE ROAD GRAPH).

Threads	Batch time	Speedup	Efficiency
1	44.39 s	1.00x	1.00
2	21.48 s	2.07x	1.03
4	10.11 s	4.39x	1.10
8	6.27 s	7.08x	0.89
16	4.56 s	9.73x	0.61

The exploratory diagnostics explain the table: the practical rule is not "A* is faster than Dijkstra." It is: *A* is faster when the heuristic is informative enough that the expansion reduction exceeds the per-expansion overhead.* Weighted A* can buy speed when the heuristic is weak but admissible: on random obstacles it reaches 20.1x for 1.8% median cost, while the perfect maze remains a zero-price control because its path is unique. On the road graph the heuristic is already tight, so reopening costs more than it saves and the closed-list gain

TABLE IX
HYPOTHESIS OUTCOMES.

	Prediction	Outcome
H1	Open: A* saves many expansions and time	confirmed
H2	Random: saves, less than open	confirmed
H3	Clustered/rooms weaker; rooms weaker than random	confirmed
H4	Maze: no expansion saving; A* slower	confirmed
H5	Long queries benefit more than short	partly rejected
H6	ALT K=8 strong; K=1 weak; plateau at 16	partly rejected
H7	`larger_g` reduces median expansions on open	rejected

is modest (7.4x to 8.8x at 1.4% cost). The OpenMP batch result shows a separate, engineering gain: independent queries scale to 9.73x throughput at 16 threads without changing which algorithm wins per query. Open grids satisfy this overwhelmingly; random, clustered and room grids satisfy it comfortably; the road graph satisfies it with ALT; the maze does not.

The result also explains why a strong but expensive heuristic can still win on the road graph. K=16 evaluates 16 landmark differences per generated vertex, but the expansion count falls enough that the net time is lower than K=8. The trade-off is graph-dependent, not a fixed property of ALT.

VII. THREATS TO VALIDITY

The main external-validity threat is the single real graph. `roadNet-CA` is a sparse, near-planar road network; the result should not be generalized to dense social graphs, directed one-way networks or weighted travel-time routing without repeating the experiment. SNAP supplies no road lengths or travel times, so the real graph is unit-cost. The grids are synthetic controls; they isolate structure but are not a sample of real maps. We symmetrise roads, so one-way streets are not modelled.

The timing is single-machine wall-clock time. The reported run was collected with the host in its lowest performance/power mode. Absolute nanoseconds are therefore mode-specific; the paired ratios are the primary result. We reran a matched subset in the host's higher performance mode: 150 grid queries at sizes 128 and 512 plus all 30 road queries, with the frozen seeds, bands and repeats. Both algorithms became about 2.24x faster, while the Dijkstra/A* ratio changed by only 1.011 on grids (95% CI 0.994–1.028, $p = 0.90$) and 0.993 on the road graph (0.967–1.020, $p = 0.21$); expansion counts were identical. We reduce noise by pinning one CPU, using one warm-up, repeating runs and reporting medians, IQRs, bootstrap intervals and paired tests. The median within-query CV is 7–10%, much smaller than the reported treatment ratios, except for the explicitly non-significant K=1 road result. A different heap, a decrease-key implementation or a different compiler could change constants, but not the mechanism: A* pays more per expansion.

The open-grid result is partly a tie-breaking effect. With `node_id`, open-grid A* still expands 54x fewer vertices than Dijkstra, but `larger_g` raises that to 222x. The main treatment was fixed in the protocol; the tie-break study is reported separately.

VIII. CONCLUSION

We compared Dijkstra and A* under a frozen paired protocol on five grid families and a real road network. A* is worth using when the heuristic captures the structure of the remaining path: on open grids it is 105.7x faster with 222.4x fewer expansions, and on a road network with ALT K=8 it is 7.62x faster with 11.54x fewer expansions (K=16: 13.52x and 22.41x). It is not worth using in a perfect maze, where it is 24% slower because Manhattan distance does not see the forced turns. Across all cases, time savings are smaller than expansion savings because A* costs 1.2–2.0x more per expansion. The right question is therefore not whether A* is better in general, but whether the heuristic is informative enough for the expansion saving to pay for the guidance. A follow-up diagnostic supports this: heuristic error predicts the speedup, and weighted A* with a closed list turns a small path-cost budget into a large gain when the heuristic is weak (random obstacles: 20.1x for 1.8% median cost). On the road graph, where the heuristic is already tight, reopening costs more than it saves; the closed-list variant improves the optimal ALT speedup from 7.4x to 8.8x at 1.4% cost. The real-graph tie-break check also confirms that `larger_g` is the safer default: the `node_id` rule expands 1.14x as many vertices and is worse on 27 of 30 queries. For batch workloads, independent queries add a further 9.73x throughput at 16 threads while preserving the per-query ranking. Finally, rerunning a matched subset in the host's higher performance mode made both algorithms about 2.24x faster but changed the Dijkstra/A* ratio by only 0.011 on grids and -0.007 on the road graph, so the ranking is not an artefact of the host power mode.

REPRODUCIBILITY

The source code, test suite, frozen protocol, raw CSV files, analysis scripts and figures are included in the project repository. The main run uses seed 20261030, 3 grid repeats and 5 road repeats, and all 3,600 runs are checked against a BFS oracle. The full experiment is reproduced with `./scripts/run_experiments.sh`; the exploratory analyses with `./scripts/run_novelty.sh`; and the OpenMP study with `./scripts/run_omp.sh`.

REFERENCES

[1] E. W. Dijkstra, "A note on two problems in connexion with graphs," *Numerische Mathematik*, vol. 1, pp. 269–271, 1959.

[2] P. E. Hart, N. J. Nilsson, and B. Raphael, "A formal basis for the heuristic determination of minimum cost paths," *IEEE Trans. Syst. Sci. Cybern.*, vol. 4, no. 2, pp. 100–107, 1968.

[3] A. V. Goldberg and C. Harrelson, "Computing the shortest path: A* search meets graph theory," in *Proc. ACM–SIAM Symp. Discrete Algorithms (SODA)*, 2005, pp. 156–165.

[4] S. Russell and P. Norvig, *Artificial Intelligence: A Modern Approach*, 4th ed. Hoboken, NJ, USA: Pearson, 2020.

[5] J. Leskovec and A. Krevl, "SNAP Datasets: Stanford large network dataset collection," 2014. [Online]. Available: https://snap.stanford.edu/data/roadNet-CA.html

[6] T. H. Cormen, C. E. Leiserson, R. L. Rivest, and C. Stein, *Introduction to Algorithms*, 4th ed. Cambridge, MA, USA: MIT Press, 2022.